

WEEK 3 — OPTIMIZED MODEL EVALUATION REPORT

1. Objective

The objective of this phase was to enhance the performance of the neural network model developed in previous stages by applying advanced optimization techniques. The focus was on improving the model's ability to detect loan defaulters, which represent the minority class in the dataset. Given the severe class imbalance, special attention was given to improving recall while maintaining a reasonable balance with precision.

2. Baseline Model Limitation

The baseline models developed in earlier experiments demonstrated high accuracy, typically around 91–92%. However, further analysis revealed that these models suffered from extremely low recall, often below 5%. This indicated that the models were heavily biased toward predicting the majority class (non-default) and failed to identify defaulters effectively.

This limitation highlighted that accuracy was not a reliable metric for evaluating performance in this problem. Instead, recall and F1 score were identified as more appropriate metrics, as they better capture the model's ability to detect critical cases.

3. Optimization Techniques Applied

To address the limitations of the baseline model, several optimization techniques were applied systematically.

3.1 Class Weighting

Class weighting was introduced to address the imbalance between default and non-default cases. By assigning higher weights to the minority class, the model was encouraged to prioritize learning patterns associated with defaulters.

```
{0: np.float64(0.5439092983356032), 1: np.float64(6.193554884189325)}
```

Impact:

- Significant improvement in recall

- Increased sensitivity to minority class
- Reduction in accuracy due to shift in prediction behavior

3.2 Dropout Regularization

Dropout layers were incorporated into the neural network to reduce overfitting. By randomly deactivating neurons during training, dropout helps the model generalize better to unseen data.

Impact:

- Improved generalization
- Reduced dependency on specific features
- No direct impact on recall, but supported overall stability

3.3 Batch Normalization

Batch normalization layers were added to stabilize and accelerate training. This technique normalizes intermediate layer outputs, reducing internal covariate shift.

Impact:

- Faster convergence
- Improved training stability
- Better gradient flow

3.4 Learning Rate Scheduling

A learning rate scheduler was used to reduce the learning rate when validation performance plateaued. This allowed finer adjustments during later stages of training.

Impact:

- Prevented oscillations in loss
- Improved convergence behavior
- Enhanced final model performance

3.5 Early Stopping

Early stopping was implemented to halt training when validation performance stopped improving. This prevented overfitting and ensured that the model retained its best weights.

Impact:

- Reduced overfitting
- Optimized training duration
- Preserved best-performing model state

3.6 Threshold Optimization

The default classification threshold of 0.5 was adjusted to explore different trade-offs between precision and recall. Multiple thresholds were evaluated to identify the optimal decision boundary.

Observed Results:

Threshold	Precision	Recall	F1 Score
0.3	0.113	0.874	0.200
0.4	0.135	0.764	0.229
0.5	0.159	0.642	0.256
0.6	0.185	0.518	0.273

Conclusion:

Threshold 0.6 was selected as it provided the best balance between precision and recall, as indicated by the highest F1 score.

4. Final Model Architecture

The optimized model consists of:

- Dense layers with sizes 256, 128, and 64
- ReLU activation functions
- Batch normalization after each major layer
- Dropout layers with a rate of 0.3
- Sigmoid output layer for binary classification

This architecture balances model complexity with generalization capability.

5. Final Model Performance

The final model was evaluated using the optimized threshold (0.6). The performance metrics are as follows:

Metric	Value
Accuracy	0.8294717330861909
Precision	0.22276879831342233
Recall	0.44692849949647534
F1 Score	0.2973335119924963


6. Performance Analysis

The optimized model demonstrates a significant improvement in recall compared to the baseline. Initially, recall was approximately 0.01, indicating almost no detection of defaulters. After applying class weighting and threshold tuning, recall increased to approximately 0.52.

This improvement confirms that the model is now capable of identifying a substantial portion of defaulters. However, this improvement comes at the cost of reduced precision, indicating an increase in false positives.

The decrease in accuracy is expected and reflects the model’s shift from majority-class bias to a more balanced classification approach.

7. Confusion Matrix Analysis



The confusion matrix reveals the distribution of predictions:

- True positives increased significantly, indicating improved detection of defaulters
- False positives also increased, contributing to lower precision
- The model prioritizes identifying defaulters, which is appropriate for risk-sensitive applications

This behavior aligns with the objective of minimizing false negatives, which are more costly in financial contexts.

8. Error Analysis

SK_ID_CURR	NAME_CONTRACT_TYPE	COD_E_GENDER	FLAG_G_OWN_CAR	FLAG_OWN_REALTY	CNT_CHI_LDR_EN	AMT_INCOME_TOTAL	AMT_CREDIT	AMT_ANNUITY	AMT_GOODS_PRICE	Actual	Predicted
220127	Cash loans	M	N	Y	0	148500.0	284400.0	18643.5	225000.0	0	1
218614	Cash loans	F	N	Y	0	112500.0	170640.0	13612.5	135000.0	0	1
394494	Revolving loans	F	N	Y	0	202500.0	225000.0	11250.0	225000.0	1	0
208491	Cash loans	M	N	Y	1	121500.0	269550.0	21294.0	225000.0	0	1
281928	Cash loans	F	Y	Y	0	180000.0	997335.0	35464.5	832500.0	1	0

A detailed examination of misclassified instances revealed several patterns:

- The model struggles with borderline cases where feature values overlap between classes
- Some non-defaulters are incorrectly classified as defaulters due to similar financial profiles
- Missing or noisy data contributes to incorrect predictions
- Certain features may lack sufficient discriminative power

These observations suggest that further improvements may require feature engineering or additional data.

9. Comparison with Baseline Model

Metric	Baseline	Optimized
Accuracy	~0.92	~0.70
Recall	~0.01	~0.52
F1 Score	Very low	0.273

The optimized model demonstrates a substantial improvement in recall, transforming it from an ineffective model to one capable of meaningful predictions. Although accuracy decreased, the model’s practical utility improved significantly.

10. Limitations

Despite improvements, the model has several limitations:

- Precision remains relatively low, indicating a high number of false positives
- Class imbalance continues to affect performance
- The model is sensitive to threshold selection
- Limited feature separability reduces classification accuracy

11. Conclusion

The application of optimization techniques significantly improved the model’s ability to detect loan defaulters. The transition from a high-accuracy but ineffective model to a balanced and

practical model demonstrates the importance of addressing class imbalance and carefully tuning decision thresholds. The final model achieves a meaningful balance between precision and recall, making it suitable for real-world applications where identifying defaulters is critical.